

PROMPT ENGINEERING

MODEL USED: ChatGpt

1.Zero shot + Role based:

NORMAL PROMPT: Write a cold email to a potential client about our analytics software.

RESULT:

Hi [Client Name],

I hope you're doing well.

I'm reaching out because many businesses today are collecting more data than ever but often struggle to turn that data into actionable insights. That's where our analytics software can help.

Our platform enables organizations to:

- Track key business metrics in real time
- Generate interactive dashboards and reports
- Identify trends and opportunities faster
- Improve decision-making through data-driven insights
- Save time by automating reporting processes

Companies using our solution have been able to gain greater visibility into their operations and make more informed strategic decisions.

I'd love to schedule a brief 15-minute call to learn about your current analytics processes and explore whether our platform could add value to your organization.

Would you be available sometime next week?

Thank you for your time, and I look forward to hearing from you.

Best regards,

[Your Name]

[Job Title]

[Company Name]

[Phone Number]

[Email Address]

PROMPT WITH ROLE:

You are a senior sales rep who has closed 200+ enterprise deals. You know that cold emails fail when they talk only about features. They work when they name a specific pain point, Write a cold email to the VP of Marketing about our analytics software.

RESULT:

Hi [First Name],

Most marketing leaders I speak with aren't struggling to collect data.

They're struggling to trust it.

Google Analytics says one thing. CRM reports say another. Paid media platforms claim credit for the same conversion. Meanwhile, your team spends hours building reports instead of improving campaigns.

That's exactly why we built [Software Name].

It gives marketing teams a single source of truth for campaign performance, customer journeys, and revenue attribution—without jumping between multiple dashboards.

Our clients use it to:

- Identify which channels actually generate pipeline and revenue
- Spot underperforming campaigns before they waste budget
- Eliminate manual reporting and spreadsheet work
- Make faster decisions with real-time insights

I'm curious—how confident are you that your current reporting accurately shows what's driving revenue?

If it's worth exploring, I'd be happy to share how other marketing teams are solving this problem in a brief 15-minute conversation.

Would next week be a bad time?

Best,

[Your Name]

[Title]

[Phone]

[Email]

2.Few shot+ Structured:

PROMPT 1:

Feature Update (Sleep Tracking) Subject: We watched you sleep. In a cool way. PulseX now tracks your sleep stages — light, deep, REM — so you can finally stop guessing why you feel like garbage at 2pm. Here's how it works: wear your band to bed. Wake up. Open the app. See your sleep score. Judge yourself accordingly. It also tracks how often you woke up, how long it took you to fall asleep, and whether your "8 hours" was actually 5 hours of tossing and 3 hours of real rest. No setup needed. The update's already live. Just open the app and tap Sleep. Fair warning — the data might hurt your feelings. [Check your sleep score →]

RESULT:

PulseX now tracks your sleep stages — Light, Deep, and REM — so you can finally stop guessing why you're running on fumes by 2 PM.

Here's how it works:

Wear your PulseX Band to bed.

Wake up.

Open the app.

See your sleep score.

Make peace with your life choices.

You'll also get insights into:

- How often you woke up during the night
- How long it took you to fall asleep
- Whether your "8 hours of sleep" was actually 5 hours of rest and 3 hours of staring at the ceiling

The best part? There's nothing to set up.

The update is already live. Just open the app and tap Sleep.

One warning: the data may be more honest than your alarm clock.

PROMPT 2:

Sale Announcement (End of Season)

Subject: 30% off. No catch. Okay, one catch. The catch is it ends Sunday. Everything in the store — bands, chargers, straps, the works — 30% off. We'd say "while stocks last" but honestly we have plenty. We just want you to feel urgency. Is it working? If you've been thinking about upgrading your strap or grabbing a second charger for the office, this is the moment. If you haven't been thinking about it, maybe start. We're not going to beg. But we're also not running this sale again until we feel like it. Code: PULSE30 at checkout. Ends Sunday 11:59pm. [Shop the sale →]

RESULT:

Here's a concise analysis of this campaign email:

Email Type

End-of-Season Sale Announcement

Objective

To drive immediate purchases by promoting a limited-time 30% discount across all products.

Target Audience

- Existing PulseX customers
- Potential customers considering a purchase
- Users interested in accessories and upgrades

Key Message

All PulseX products are available at **30% off** for a limited time using the code **PULSE30**, with the sale ending Sunday at 11:59 PM.

Features Highlighted

- 30% discount storewide
- Includes bands, chargers, and straps
- Simple discount code: PULSE30
- Limited-time offer

Tone of Voice

- Humorous
- Casual
- Self-aware
- Conversational

Call to Action (CTA)

"Shop the sale →"

Marketing Techniques Used

- Urgency ("Ends Sunday")
- Limited-time offer
- Humor and transparency
- Fear of missing out (FOMO)
- Direct and simple messaging

Why This Email Is Effective

- The subject line immediately communicates the discount.
- Creates urgency without sounding overly salesy.
- Uses humor to make the promotion memorable.
- Clearly states the offer and expiration date.

- Encourages both planned and impulse purchases.

AIDA Analysis

Element Example

Attention "30% off. No catch. Okay, one catch."

Interest Storewide discount on all products.

Desire Opportunity to upgrade accessories at a lower price.

Action Use code **PULSE30** and click "**Shop the sale →**"

Customer Psychology Used

- **Scarcity:** Sale ends Sunday.
- **FOMO:** Customers may miss a significant discount.
- **Humor:** Makes the brand feel relatable and trustworthy.
- **Convenience:** One simple code for all products.
-

CRAFT PROMPT:

<Context> I am currently designing a production grade pipeline using kafka + DBT + Snowflake and terraform integration. My team is full of high quality engineers, and hence I want to make sure, that the code scripted by me too stands out and abides by all the software engineering principles. I want this code to be scalable and optimized in the best possible way.

Role: Act as Senior software engineer with 2 decades of experience, who has worked on high quality enterprise-level scalable applications and is adept with data engineering knowledge and impeccable programming skills. You have worked on a variety of streams like data engineering, machine learning, database engineering. You have unique way for explaining complex concepts in a simpler manner, and also reviewing code in the best possible manner and suggesting the most potent additions to the code which enhances the application quality.

<Action>: I want you to review the code and suggest me where I the code be engineered in a better way, where can it be optimized, which part of the code seems unnecessary and where can i improve better. Which are the strengths and weaknesses in the code, and also suggest me the write versions of the code, after suggesting the changes.

<Format: > OUtput should be strictly in markdown format. Only if required, perform comparative analysis in tabular format.

Constraints: The code review should be strictly with the code uploaded. Avoid explaining too technically, rather it should be simpler to follow and implement. Be precise and verbose at the same time.

RESULT:

Review Structure

1. Executive Summary

- Overall code quality score
- Production readiness assessment
- Scalability assessment
- Maintainability assessment

2. Strengths

- What is implemented well
- Good engineering decisions
- Reusable patterns
- Performance-conscious implementations

3. Issues & Improvement Areas

For each issue:

Issue

- What is wrong
- Why it matters

Impact

- Performance
- Scalability
- Reliability
- Maintainability

Recommendation

- Simpler explanation
- Best-practice approach

Improved Code

- Refactored version of the specific section

4. Data Engineering Review

Specific review for:

- Kafka design
- Consumer groups
- Partition strategy
- Event schema design
- Idempotency
- Retry handling
- Dead-letter queues
- Snowflake optimization
- DBT modeling
- Incremental models
- Terraform structure
- CI/CD readiness

5. Scalability Review

Area Current State Risk Recommended Change

6. Production Readiness Checklist

Review of:

- Logging
- Monitoring
- Alerting
- Error handling
- Security
- Secrets management
- Observability
- Testing strategy
- Disaster recovery

7. Code Smells

- Unnecessary code
- Redundant logic
- Hidden bugs

- Future maintenance risks

8. Enterprise-Level Enhancements

Suggestions that would make the code stand out in a team of senior engineers.

9. Final Refactored Version

- Cleaned-up implementation
- Better naming
- Better structure
- Improved performance
- Production-grade recommendations

MEETING TRANSCRIPT:

Weekly AI Pipeline Sync — Document Intelligence Platform

Date	Thursday, March 12, 2026 — 10:30 AM IST
Duration	47 minutes
Location	Google Meet (recorded)
Attendees	Priya Sharma (VP of Product), Arjun Mehta (ML Engineering Lead), Sneha Iyer (Data Engineering Manager), Rohit Desai (Senior ML Engineer), Kavya Nair (Product Manager)
Absent	Nikhil Rao (CTO) — at board meeting

TRANSCRIPT BEGINS

[00:00]

Priya Sharma: Alright, let’s get started. Nikhil’s at the board meeting so he won’t be joining, but he wants an update afterward, especially on the timeline. So let’s make sure we’re clear on where things stand. Arjun, do you want to kick off with the model update?

[00:32]

Arjun Mehta: Sure. So the good news is the extraction model is performing significantly better after the retraining last week. We switched from the fine-tuned GPT-3.5 to a Mistral 7B base

with LoRA adapters, and the field-level accuracy on invoices went from 78% to 91%. That's on our internal test set of 2,200 documents.

Priya Sharma: That's a big jump. Is that across all document types or just invoices?

Arjun Mehta: Just invoices for now. Purchase orders are at about 84%, which is up from 72%, so also a solid improvement. But contracts are still a problem. We're hovering around 67% on the contract extraction, and honestly I'm not sure fine-tuning alone is going to get us where we need to be on those.

[01:45]

Rohit Desai: Yeah, the contract issue is structural. Contracts have way more variability in layout — different firms, different jurisdictions, different formatting conventions. We've been talking about a two-stage approach where we first classify the contract type and then route it to a specialized extraction head, but that adds a lot of complexity to the pipeline.

Priya Sharma: How much complexity are we talking? Another two weeks? A month?

Rohit Desai: Honestly, if we want to do it properly, probably three to four weeks of engineering time. And that's just the model work — Sneha's team would need to build out the routing logic on the data pipeline side too.

[02:48]

Sneha Iyer: Hold on, let me push back on that a little. The routing logic isn't the hard part. We can have that up in a few days, it's basically a classifier output feeding into an Airflow branch operator. The hard part is the training data. Do we even have enough labeled contracts to train specialized extraction heads? Last I checked we had maybe 400 labeled contracts total, and that's across all types.

Arjun Mehta: We have 430, to be exact. And if we split by type — NDAs, service agreements, licensing agreements — the smallest bucket is licensing with only 67 examples. That's not enough to fine-tune anything reliable.

Priya Sharma: So what's the alternative? We can't ship with 67% accuracy on contracts. That's the one document type our enterprise clients actually care about the most.

[04:12]

Arjun Mehta: I think the realistic options are: one, we invest in a labeling sprint and get the training data up to at least 200 per contract type before we attempt the two-stage approach. Two, we use a prompted approach for contracts instead of fine-tuning — basically long-context extraction with Claude or GPT-4 using carefully engineered prompts. Three, we descope contracts from the V1 launch and add them in V1.1.

Kavya Nair: Option three is not going to fly with the sales team. I've already got two enterprise prospects — Reliance Digital and Tata Communications — who specifically asked about contract extraction in their last calls. If we pull that from V1, we might lose both deals.

Priya Sharma: Agreed, descoping is a last resort. What about option two — the prompted approach? Arjun, is that viable?

[05:30]

Arjun Mehta: It's viable but it changes our cost structure. Right now the fine-tuned Mistral model runs on our own GPU cluster — inference cost is basically zero after the initial hardware investment. If we switch contracts to a prompted approach with an external API, we're looking at roughly 12 to 15 paisa per page, which adds up fast for contract-heavy clients.

Rohit Desai: It's also slower. Fine-tuned model does extraction in about 1.2 seconds per page. A prompted approach with Claude is more like 6 to 8 seconds depending on contract length. For a 40-page contract, that's over five minutes.

Kavya Nair: Is five minutes a dealbreaker though? These aren't real-time use cases. A client uploads a contract and gets the extraction back in five minutes — that's still way faster than their current process, which is someone manually reading the thing for an hour.

[06:52]

Sneha Iyer: I'm less worried about the latency and more worried about the reliability. If we're calling an external API for a core feature, we need proper retry logic, fallback handling, rate limiting — all of that. And we need to figure out the data residency question. Some of these enterprise clients have strict requirements about where their data gets processed. Are we okay sending their contracts to Anthropic's API?

Priya Sharma: That's a good point. Kavya, do we know the data residency requirements for Reliance and Tata?

Kavya Nair: Reliance definitely has a data residency clause — everything has to stay within India. I'd need to check the specifics for Tata. Let me follow up with their procurement team and get clarity this week.

Priya Sharma: Okay, please do that by Monday. We can't make a technical decision without knowing the constraints.

[08:15]

Priya Sharma: Alright, here's what I'm thinking. Let's pursue option one and two in parallel. Arjun, can you and Rohit start the labeling sprint for contracts? Get the annotation team to prioritize contract labeling — I want at least 200 labeled examples per contract type within two weeks.

Arjun Mehta: That's aggressive. The annotation team is currently finishing the invoice validation batch. We'd need to pull them off that.

Priya Sharma: The invoice model is already at 91%. We can pause the invoice annotations. Contracts are the bottleneck.

Arjun Mehta: Fair enough. I'll talk to Meera on the annotation team today and get them redirected.

[09:35]

Priya Sharma: Good. And simultaneously, Rohit, can you prototype the prompted approach? I want to see what accuracy we can get with a well-engineered prompt on the 430 contracts we already have. Give me a benchmark by end of next week.

Rohit Desai: I can do the benchmark, but end of next week is tight. I'm also supposed to be fixing the table extraction bug that's been blocking the QA team.

Sneha Iyer: That table extraction bug is actually causing issues on my side too. The downstream ETL pipeline breaks every time it encounters a document with malformed table outputs. We've had 14 pipeline failures this week alone because of it.

Priya Sharma: Okay, how severe is the table bug? Is it a 'things crash' problem or a 'things are slightly wrong' problem?

Sneha Iyer: Things crash. The pipeline fails, the document gets sent to a dead letter queue, and we have to manually reprocess. It's costing us about two hours of engineering time per day just babysitting it.

[11:20]

Priya Sharma: That's not acceptable. Rohit, fix the table bug first. That's priority one. Then start the contract prompt benchmark. Can you have the bug fixed by Monday?

Rohit Desai: I think so, yeah. I've already identified the root cause — it's a parsing issue with nested tables. The model outputs valid JSON but the parser doesn't handle the nesting correctly. Should be a day of work, maybe two.

Priya Sharma: Good. Fix it by Monday, then start the contract benchmark. Get me results by the Friday after — March 27th.

Rohit Desai: Got it.

[12:30]

Kavya Nair: One more thing on the timeline. Our V1 target launch date is April 18th. That's five weeks from now. Are we still tracking for that? Because if the contract approach takes four more weeks of engineering after the benchmark results come in, we're already past the launch date.

Arjun Mehta: Honestly? I think April 18th is at risk. Even if the prompted approach works, integrating it properly with fallback handling, monitoring, and cost tracking is at least two weeks of work. And that's after we know which approach we're going with.

Priya Sharma: I hear you, but I'm not ready to move the date yet. Let's see where we are after the benchmark results on March 27th. If we need to adjust, we'll adjust then. But I don't want us to preemptively relax the timeline. Sometimes deadlines make magic happen.

Kavya Nair: Fair, but can we at least flag it to Nikhil? He's going to ask.

Priya Sharma: I'll talk to Nikhil. Let me frame it properly — we're on track for invoices and purchase orders, contracts are the open question, and we'll have clarity by end of month.

[14:45]

Priya Sharma: Okay, let's move to the data pipeline update. Sneha, where are we on the ingestion service?

Sneha Iyer: Good progress, actually. The async ingestion pipeline is fully built and deployed to staging. We're processing about 1,200 documents per hour in testing, which is well above our target of 800. The S3-to-processing queue is stable, Airflow orchestration is clean, and we've got proper dead letter queues for failed documents.

Sneha Iyer: The one outstanding issue — apart from the table bug that Rohit mentioned — is the OCR pre-processing for scanned documents. We're using Tesseract right now, and the quality is... not great on low-resolution scans. About 22% of our test documents are scanned, and for those, the extraction accuracy drops by 15 to 20 points because the OCR introduces errors before the model even sees the text.

[16:30]

Arjun Mehta: Have you looked at Azure Document Intelligence? Their OCR is significantly better than Tesseract, especially on degraded scans. We used it at my last company and it was night and day.

Sneha Iyer: I've looked at it. The accuracy is better, no question. But it's another external API dependency, another cost line, and another data residency conversation.

Kavya Nair: Also, sorry to keep bringing this up, but Azure has an India region, right? That might actually solve the data residency issue for both the OCR and the contract extraction if we go with a prompted approach through Azure OpenAI.

Arjun Mehta: Azure OpenAI doesn't have the latest Claude models though. If we're going prompted, Claude is performing noticeably better than GPT-4 on our extraction benchmarks. Like, 6 to 8 points better on contracts specifically.

[18:10]

Priya Sharma: Okay, we're going in circles a bit. Let me try to untangle this. Sneha, can you run a quick comparison of Tesseract versus Azure Document Intelligence this week? Use 50 of the worst scanned documents. I just need to know if the accuracy improvement justifies adding another vendor dependency.

Sneha Iyer: Yeah, I can have that done by Wednesday.

Priya Sharma: Great. And the data residency question — Kavya is going to find out the requirements from Reliance and Tata by Monday. Once we have that, a lot of these API decisions become clearer. Let's not go back and forth on Azure versus Anthropic until we actually know what's allowed.

Kavya Nair: Makes sense.

[19:45]

Priya Sharma: Anything else before we wrap? We're almost at time.

Rohit Desai: One quick thing — the GPU cluster is at 87% utilization. When we scale the annotation team and start running more training jobs for the contract models, we're going to hit capacity. Should I file a request for additional GPU allocation now, or wait?

Arjun Mehta: File it now. The procurement process takes two weeks minimum. Better to have it in the queue and cancel than to need it and not have it.

Priya Sharma: Agreed. Rohit, submit the GPU request today. Estimate what we'd need for the contract training runs and add a 30% buffer.

Rohit Desai: Will do.

[20:55]

Kavya Nair: Oh wait, one more thing. The product demo for the Reliance team is scheduled for April 3rd. That's three weeks out. What are we showing them? Because if contract extraction isn't ready, we need to frame the demo carefully.

Priya Sharma: Good catch. Let's show invoices and purchase orders — those are solid. For contracts, we can show a "coming soon" preview or a limited demo with NDAs only, since those are our strongest contract subtype. Kavya, can you draft a demo flow by end of next week and share it with the team?

Kavya Nair: Sure, I'll put together a demo script and share it by Friday the 20th.

[22:10]

Priya Sharma: Alright. Let me do a quick recap. Arjun — redirect the annotation team to contract labeling, target 200 per type in two weeks. Rohit — fix the table extraction bug by Monday, then start the prompted contract benchmark, results by March 27th. Also submit the GPU request today. Sneha — run the Tesseract versus Azure OCR comparison by Wednesday. Kavya — get data residency requirements from Reliance and Tata by Monday, and draft the demo flow by the 20th. I'll update Nikhil on the timeline situation. Everyone clear?

Arjun Mehta: Clear.

Sneha Iyer: Got it.

Rohit Desai: Yep.

Kavya Nair: All good.

Priya Sharma: Great. Good progress on the model side — 91% on invoices is a real milestone. Let's keep that momentum going. Talk next week.

TRANSCRIPT ENDS — 22:48

MEETING SUMMARY:

Weekly AI Pipeline Sync — Meeting Summary

Date: March 12, 2026
Duration: 47 Minutes
Attendees: Priya Sharma, Arjun Mehta, Sneha Iyer, Rohit Desai, Kavya Nair
Absent: Nikhil Rao

Executive Summary

The team reviewed progress on the Document Intelligence Platform ahead of the planned April 18 V1 launch.

Key highlights:

- Invoice extraction accuracy improved significantly from 78% to 91% after migrating to Mistral 7B + LoRA adapters.
- Purchase order extraction improved from 72% to 84%.
- Contract extraction remains the primary risk area at 67% accuracy.
- Async ingestion pipeline exceeded performance targets, processing 1,200 documents/hour versus the target of 800 documents/hour.
- A critical table extraction bug is causing pipeline failures and requires immediate attention.
- OCR quality for scanned documents remains a challenge.
- GPU cluster utilization has reached 87%, creating potential capacity constraints for future model training.

Model Performance Update

Document Type	Previous Accuracy	Current Accuracy
Invoices	78%	91%
Purchase Orders	72%	84%
Contracts	N/A	67%

Contract Extraction Challenges

Current challenges include:

- High variability in contract layouts.

- Limited training data.
- Insufficient labeled examples for specialized fine-tuning.

Current labeled contract inventory:

Contract Type	Labeled Examples
---------------	------------------

Total Contracts	430
-----------------	-----

Licensing Agreements	67 (smallest category)
----------------------	------------------------

Contract Extraction Options Discussed

Option 1: Labeling Sprint + Two-Stage Architecture

Approach:

1. Classify contract type.
2. Route to specialized extraction model.

Pros:

- Better long-term accuracy.
- Fully owned infrastructure.

Cons:

- Requires significant additional labeled data.
- Estimated 3–4 weeks of engineering effort.

Option 2: Prompted Extraction

Approach:

- Use Claude or GPT-4 with prompt engineering.

Pros:

- Faster path to improved accuracy.
- No immediate need for large training datasets.

Cons:

- Increased inference costs.
- Higher latency.
- Data residency concerns.

- Additional operational complexity.

Option 3: Remove Contracts from V1

Outcome:

- Rejected as a last-resort option due to enterprise customer demand.
-

Critical Engineering Issues

Table Extraction Bug

Impact:

- 14 pipeline failures during the week.
- Documents routed to DLQ.
- Approximately 2 engineering hours/day spent on manual recovery.

Root Cause:

- Nested table parsing issue.
- Model outputs valid JSON.
- Parser fails to correctly process nested structures.

Priority:

Highest Priority

Target

Monday

Fix

Date:

Data Pipeline Update

Ingestion Service

Status: Healthy

Achievements:

- Async ingestion pipeline deployed to staging.
- Stable S3 → Queue → Processing architecture.
- Airflow orchestration functioning correctly.
- Dead-letter queue implementation completed.

Performance:

Metric	Target	Actual
Documents/hour	800	1,200

OCR Discussion

Current OCR:

- Tesseract

Issue:

- Poor performance on low-resolution scans.
- Approximately 22% of documents are scanned.
- OCR errors reduce extraction accuracy by 15–20 percentage points.

Alternative Evaluated:

- Azure Document Intelligence

Considerations:

- Better OCR accuracy.
 - Additional vendor dependency.
 - Additional cost.
 - Data residency implications.
-

Data Residency Concerns

Potential enterprise customers:

- Reliance Digital
- Tata Communications

Discussion Points:

- Reliance reportedly requires India-only data processing.
- Need formal confirmation before selecting external AI/OCR vendors.

Decision:

- Validate customer requirements before making architecture decisions regarding:
 - Anthropic
 - Azure OpenAI

- Azure Document Intelligence
-

Infrastructure Update

GPU Capacity

Current utilization:

87%

Risk:

- Contract model training could exceed available capacity.

Decision:

- Submit additional GPU allocation request immediately.
 - Include estimated requirements plus 30% buffer.
-

Timeline Assessment

Current V1 Launch Target:

April 18, 2026

Risk Assessment:

- Invoice and Purchase Order workflows remain on track.
 - Contract extraction remains the primary schedule risk.
 - Final decision deferred until benchmark results are available on March 27.
-

Customer Demo Planning

Customer:

- Reliance Digital

Demo Date:

April 3, 2026

Planned Scope:

Demonstrate

- Invoice extraction
- Purchase order extraction

Potential Preview

- NDA extraction workflow

Not Guaranteed

- Full contract extraction capability
-

Action Items

Owner	Action Item	Due Date
Arjun	Redirect annotation team to contract labeling	Immediate
Arjun	Reach 200 labeled examples per contract type	Within 2 weeks
Rohit	Fix table extraction bug	Monday
Rohit	Run prompted contract benchmark	March 27
Rohit	Submit GPU capacity request	Today
Sneha	Compare Tesseract vs Azure OCR using 50 poor-quality scans	Wednesday
Kavya	Confirm data residency requirements with Reliance and Tata	Monday
Kavya	Draft customer demo flow	March 20
Priya	Update CTO on timeline and contract extraction risk	Ongoing

Key Risks

1. Contract extraction accuracy remains below acceptable levels.
 2. Limited labeled contract dataset.
 3. Data residency constraints may restrict vendor choices.
 4. OCR quality issues affect scanned document accuracy.
 5. GPU capacity nearing limits.
 6. April 18 launch date is at risk if contract solution requires significant additional engineering.
-

Decisions Made

- Prioritize contract labeling over invoice annotation work.

- Fix table extraction bug before contract benchmark work.
- Pursue labeling sprint and prompted extraction approach in parallel.
- Delay external API decisions until data residency requirements are confirmed.
- Submit GPU expansion request immediately.
- Maintain April 18 launch target until benchmark results are reviewed on March 27.